

CSE231 - Operating Systems

Assignment 3

Exercise 3.1

Name: Abhik S Basu
Roll Number: 2020165
Section: A
Branch: CSE

Task: Modifying vruntime of Linux CFS Scheduler

To approach this question, we need to first understand how exactly the vruntime of a particular process is calculated and where exactly is it calculated. Following that we will establish a sys call:

They are as follows:

`sched.h`: /home/kern/build/linux-5.14.7/include/linux/sched.h

`fair.c`: build/linux-5.14.7/kernel/sched/fair.c

`core.c`: /home/kern/build/linux-5.14.7/kernel/sched/core.c

Here are the changes that we have made to the following files:

- 1) In `sched.h` we find the `sched_entity` structure and add an attribute to it which will store the amount of delay that has to be given to the vruntime of the process.
- 2) In `fair.c`, changes are made to the `update_curr()` function as well . In this we add the newly added delay attribute when we are computing the `curr_vruntime` of the process there we add the delay attribute.

3) In `core.c` , we initialise all the members of `sched_entity` to zero. So we just add the `delay` attribute and set that to zero as well.

Now after making these changes to the kernel we now add the definition of the syscall in the syscall table and further define it in the `sys.c` document.

In the syscall, we take in two parameters which are the ID of the process whose `vruntime` has to be increased. And the other parameter is the delay time which is taken by us in millisecond. We convert the millisecond to nanoseconds in our syscall by multiplying it by $1e6$. Here we also initialise the structure `pid_struct` since we will assign it the PID on which the delay time has to be added. After doing error handling on the `vruntime` and `pid` we do the multiplication, after that we assign the struct the PID and then using `pid_task()` function we assign the `task_struct` the `pid_struct`. Then we change the `vruntime` of this process by adding the delay, the delay has been computed and assigned to the newly added attribute. This is basically what the syscall is doing. I am also printing the newly delayed time in order to show it on the system log that how the `vruntime` has been affected for each of the processes. This will tell us the scale of exponents by which `vruntime` has been affected.

Following that we execute the `make` and then we will see the changes.

Now I will explain how the testing works:

We have created two programmes `test` and `test2`. In '`test`' we first compute the time the process runs for and then in '`test2`' we compute the time the process runs for when the syscall is invoked. We will now notice the difference in the execution time. Since the process is getting changed in the `update_curr()` so we will see the difference in the execution time.

We can test all of this by running the command as follows:

make

```
all: t1 t2 run

t1: test.c
    gcc test.c -lgomp -o test

t2: test2.c
    gcc test2.c -lgomp -o test2

clean:
    rm test test2

run:
    ./test && ./test2
```

```
[kern@artixcse231 ~]$ make
gcc test.c -lgomp -o test
gcc test2.c -lgomp -o test2
./test && ./test2
Nice value of process is : 5
Timeslice for non Sys Call process is: 0 seconds, 16666665 nanoseconds
Finished executing for non Sys Call in 3.895953 seconds
Nice value of process is : 5
Timeslice for Sys Call process is: 0 seconds, 16666665 nanoseconds
Finished executing for Sys Call in 3.948575 seconds
[kern@artixcse231 ~]$
```

```
[kern@artixcse231 ~]$ make
gcc test.c -lgomp -o test
gcc test2.c -lgomp -o test2
./test && ./test2
Nice value of process is : 5
Timeslice for non Sys Call process is: 0 seconds, 16666665 nanoseconds
Finished executing for non Sys Call in 3.996974 seconds
Nice value of process is : 5
Timeslice for Sys Call process is: 0 seconds, 16666665 nanoseconds
Finished executing for Sys Call in 4.108231 seconds
```

We can check the kernel log, by using the command as follows:

```
sudo dmesg | tail
```

The output for that will be as follows:

```
vruntime of PID: 382 is now 8111640145298071354
PID Delay : -> 10000000000000000
vruntime of PID: 383 is now 8800647191637014406
PID Delay : -> 100000000000
vruntime of PID: 383 is now 186471204942
PID Delay : -> 186471306972
```

We can notice that the difference or increase in the vruntime will always be scaled to the amount of digits in the delay. We can compare this to previous runtime as well by doing another `printk()` in our sys call definition and it will work. The difference b/w the vruntime will be of the form $1e(x-y)$, x will be digits in bigger delay and y is digits in smaller. And the difference will approximately be equal to $1e(x-y)$.